

Predict Podcast Listening Time



ICS 435 Final Project
Christian Dela Cruz, Kyler Okuma, Sean Flynn

Background

Our Goal: Predict listening time of a podcast episode.

We were provided with a training dataset containing episode information and a target variable, `Listening_Time_minutes`.

Our objective was to build predictive models to estimate listening time as accurately as possible, measured by Root Mean Squared Error (RMSE) on a separate test set.

Evaluation Metric: Root Mean Squared Error (RMSE)

RMSE was used to measure model performance. RMSE is a standard metric for regression tasks that heavily penalizes larger errors.

$$\text{RMSE} = \left(\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 \right)^{\frac{1}{2}}$$

Novelty

- New application area
 - Most ML applications focus on **classification** or **traditional regression**
 - **Predicting user engagement in audio content** has not been explored
- Comparison of ML models
 - Regression
 - MLP
 - LightGBM
 - XGBoost
 - CatBoost
- Simulated real-world ML workflow
 - Feature extraction
 - Tuning
 - Evaluation
 - Validation

Models

- Regression
 - Light
 - Ridge
 - Lasso
- MLP
 - w/ SGD
 - w/ Adam
 - w/ RMSprop
 - w/ Adagrad
 - Ensemble of 5 MLPs
 - Final Optimized MLP (Optuna Parameters)
- Optimized LightGBM
- Stacked Ensemble (All Models Above)
- XGBoost
- CatBoost

Regression

Overview:

- Used Linear Regression, Ridge Regression, and Lasso Regression.
- Assumed a linear relationship between features and listening time.
- Prioritized interpretability over flexibility.

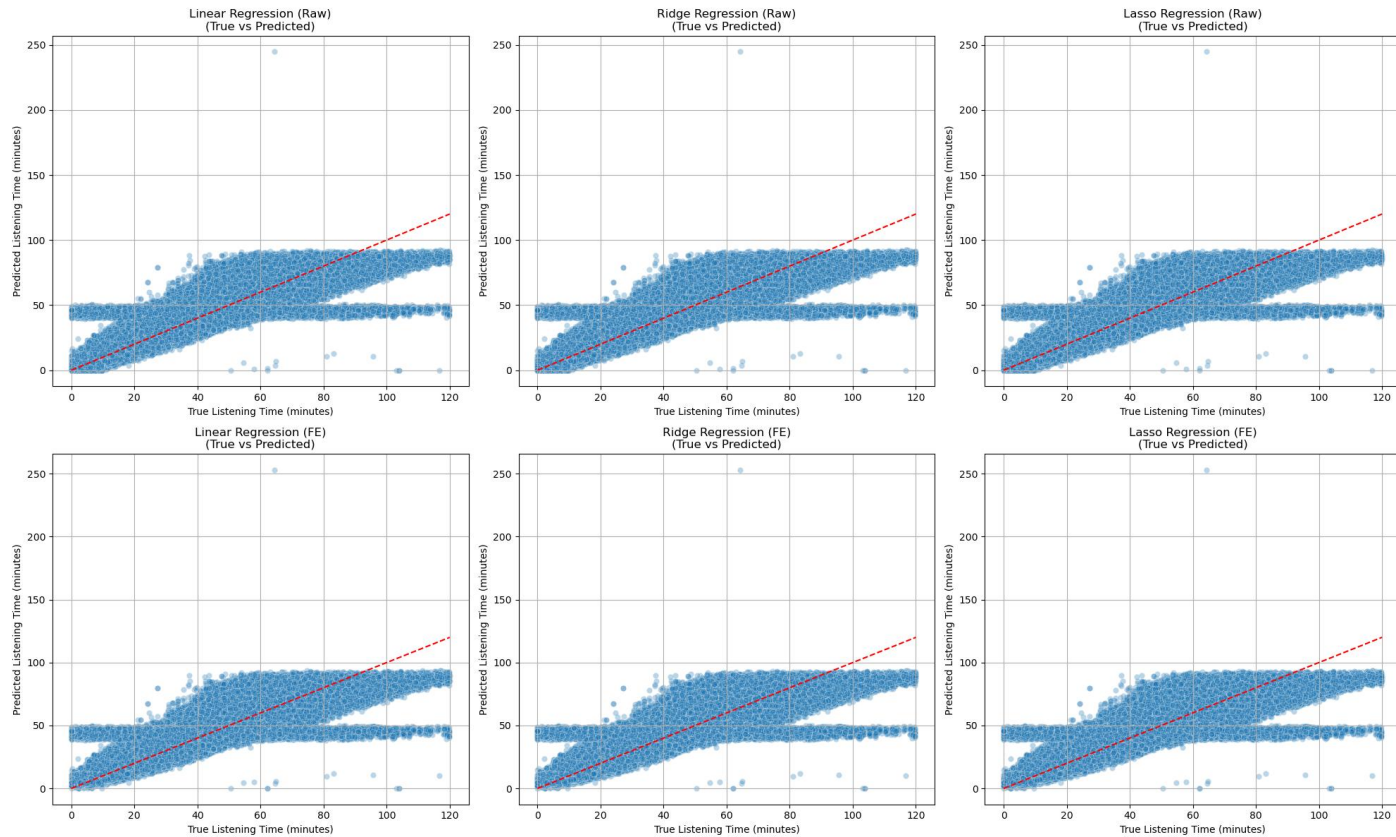
Training Approach:

- Trained separately on RAW and FE datasets.
- GridSearchCV used to tune regularization strength (alpha).
- Categorical features one-hot encoded; missing values imputed.


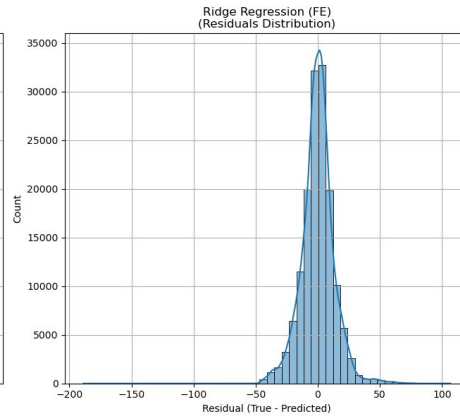
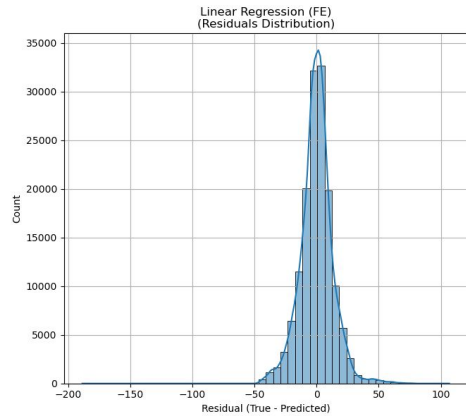
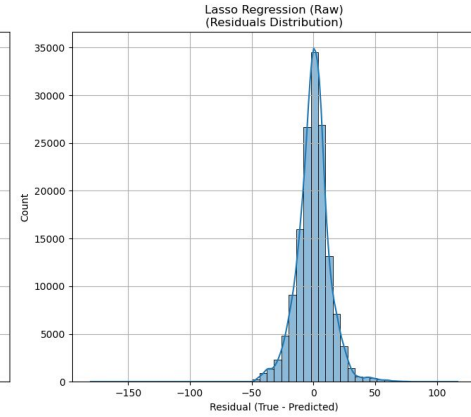
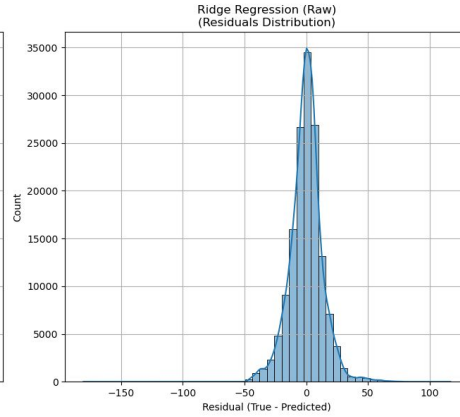

Results:

- Linear Regression (FE): RMSE = 13.296, MAE = 9.742, $R^2 = 0.7598$
- Ridge Regression (FE): RMSE = 13.295, MAE = 9.741, $R^2 = 0.7598$
- Lasso Regression (FE): RMSE = 13.295, MAE = 9.741, $R^2 = 0.7598$

Results - Linear Regression



Results - Linear Regression



MLP

Overview:

- Implemented Multi-Layer Perceptron (MLP) neural networks.
- Aimed to capture nonlinear feature interactions.
- Feedforward architecture capable of approximating complex patterns.

Training Approach:

- Trained on FE dataset with standardized features.
- Tested optimizers: SGD, Adam, RMSprop, Adagrad.
- Adagrad optimizer selected based on best validation results.
- Hyperparameters optimized using Optuna.
- Early stopping used to prevent overfitting.

Results:

- Best Single MLP: RMSE = 13.204, MAE = 9.628, $R^2 = 0.7631$
- Ensemble of 5 MLPs: RMSE = 13.271, MAE = 9.696, $R^2 = 0.7606$

Ensemble

Overview:

- Combined outputs of multiple strong models.
- Two ensemble strategies:
 - **Tree-Based Ensemble:** Averaged optimized LightGBM, XGBoost, and CatBoost models.
 - **Full Stacked Ensemble:** Stacked outputs from:
 - Linear Regression (FE)
 - Ridge Regression (FE)
 - Lasso Regression (FE)
 - Best Optuna-tuned MLP
 - Ensemble of 5 MLPs (Adagrad-trained)
 - Best Optuna-tuned LightGBM

Training Approach:

- Optimized each base model separately using GridSearchCV or Optuna.
- Averaged tree predictions for the Tree-Based Ensemble.
- Trained Ridge Regression as meta-learner for Full Stacked Ensemble.

Results:

- Tree-Based Ensemble: RMSE = 12.696, MAE = 9.248, R^2 = 0.7809
- Full Stacked Ensemble: RMSE = 12.931, MAE = 9.382, R^2 = 0.7728

Categorical Boosting

Model Overview

- Common boosting method typically used to handle categorical features
- Known for quick speed and accuracy, as it does not require encoding of categorical features




Image Credit:

https://www.researchgate.net/figure/Categorical-boosting-CatBoost-algorithm_fig5_380847150

Categorical Boosting (Cont'd)

Training Steps

- 1) Train with default parameters
- 2) Perform hyperparameter tuning with optuna
- 3) 30 trial optimization
- 4) Train with “optimal” hyperparameters
- 5) Predict on new data

```
best_catboost_params = {  
    'learning_rate': 0.053156533135846354,  
    'depth': 11,  
    'iterations': 735,  
    'l2_leaf_reg': 0.5537806027515884,  
    'bagging_temperature': 0.966594578770241,  
    'colsample_bylevel': 0.7327316098447535,  
    'border_count': 136,  
    'random_state': 42,  
    'verbose': False  
}
```

*Optimal hyperparameters trained
on model*

Categorical Boosting (Cont'd)

Results

- Low RMSE score indicates good performance
- Performed worse than other tree-based models
 - Model's full strength may not have been utilized, as not *all* of the features were categorical
 - Still followed the common trend throughout the project, with tree-based models performing the best

Final Optimized CatBoost Results:

```
{'Model': 'Final Optimized CatBoost (Optuna Hyperparams)', 'RMSE': 12.997247254066897, 'MAE': 9.454901086873276, 'R2': 0.7704238101092351}
```

Final results

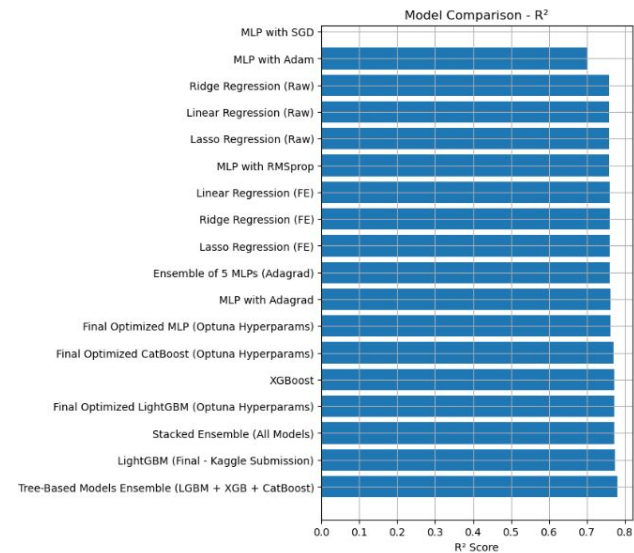
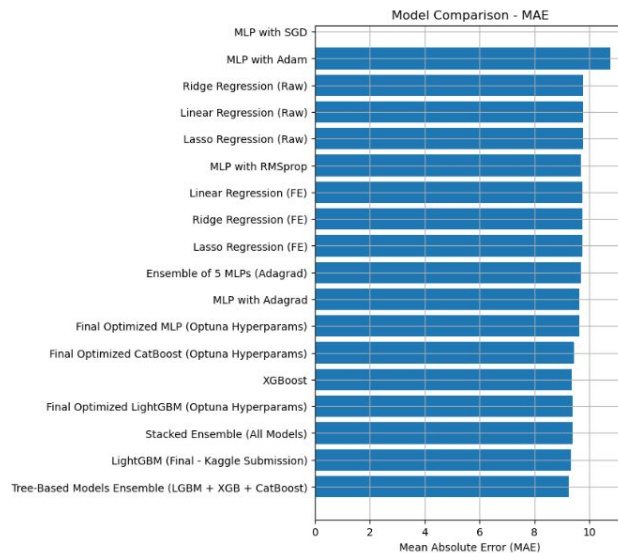
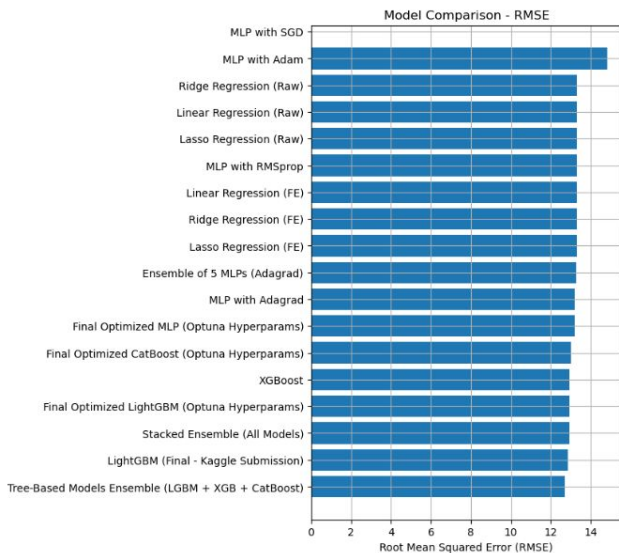
Results - XGBoost Model

```
[0]      validation_0-rmse:26.09995
[100]    validation_0-rmse:13.09422
[200]    validation_0-rmse:13.05328
[300]    validation_0-rmse:13.02910
[400]    validation_0-rmse:13.01441
[500]    validation_0-rmse:13.00028
[600]    validation_0-rmse:12.99016
[700]    validation_0-rmse:12.98256
[800]    validation_0-rmse:12.97533
[900]    validation_0-rmse:12.96928
[1000]   validation_0-rmse:12.96210
[1100]   validation_0-rmse:12.95727
[1200]   validation_0-rmse:12.95259
[1300]   validation_0-rmse:12.94992
[1400]   validation_0-rmse:12.94535
[1500]   validation_0-rmse:12.94351
[1600]   validation_0-rmse:12.94176
[1700]   validation_0-rmse:12.93899
[1800]   validation_0-rmse:12.93652
[1900]   validation_0-rmse:12.93521
[1999]   validation_0-rmse:12.93239
✅ XGBoost Validation RMSE: 12.93237
```

Results - LightGBM Model

```
[100] training's rmse: 13.001 valid_1's rmse: 13.0542
[200] training's rmse: 12.8623 valid_1's rmse: 13.016
[300] training's rmse: 12.7532 valid_1's rmse: 12.9914
[400] training's rmse: 12.6462 valid_1's rmse: 12.9731
[500] training's rmse: 12.5645 valid_1's rmse: 12.9577
[600] training's rmse: 12.4733 valid_1's rmse: 12.9443
[700] training's rmse: 12.3918 valid_1's rmse: 12.9334
[800] training's rmse: 12.3185 valid_1's rmse: 12.9239
[900] training's rmse: 12.2356 valid_1's rmse: 12.9139
[1000] training's rmse: 12.1584 valid_1's rmse: 12.9049
[1100] training's rmse: 12.0934 valid_1's rmse: 12.899
[1200] training's rmse: 12.0259 valid_1's rmse: 12.895
[1300] training's rmse: 11.9525 valid_1's rmse: 12.886
[1400] training's rmse: 11.8835 valid_1's rmse: 12.8802
[1500] training's rmse: 11.8121 valid_1's rmse: 12.8766
[1600] training's rmse: 11.748 valid_1's rmse: 12.8733
[1700] training's rmse: 11.6848 valid_1's rmse: 12.8668
[1800] training's rmse: 11.6277 valid_1's rmse: 12.8632
[1900] training's rmse: 11.5724 valid_1's rmse: 12.8582
[2000] training's rmse: 11.5165 valid_1's rmse: 12.8558
Did not meet early stopping. Best iteration is:
[2000] training's rmse: 11.5165 valid_1's rmse: 12.8558
✔ LightGBM Tuned Validation RMSE: 12.85577
```

Results - Model Comparison


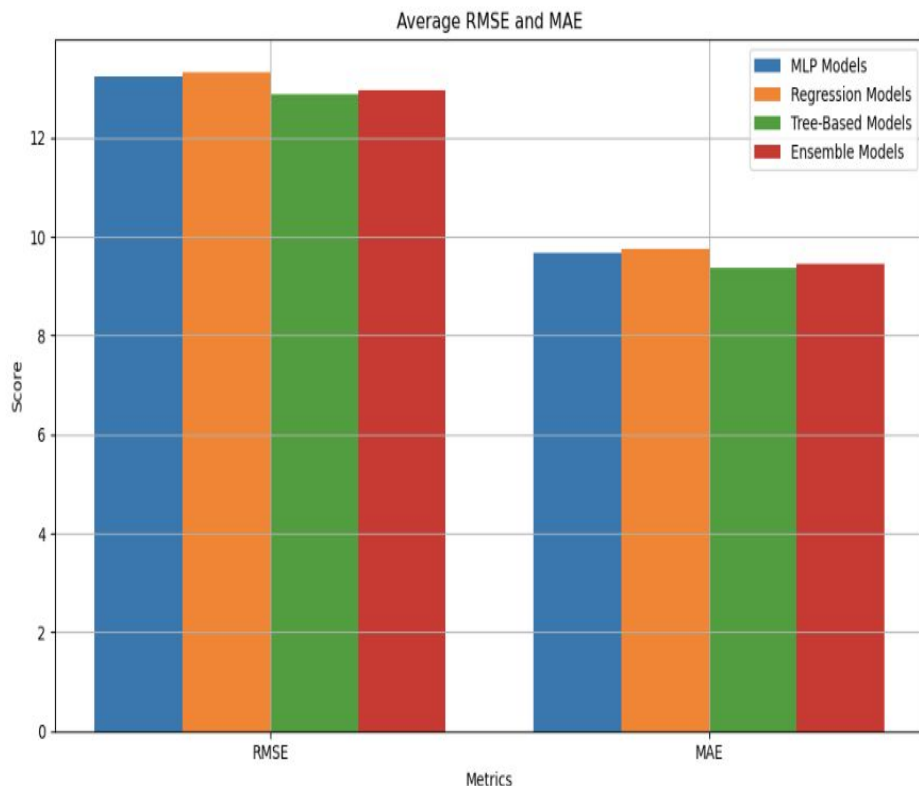


Results

	Model	RMSE	MAE	R2	Type
0	Linear Regression (Raw)	13.319984	9.764266	0.758881	Linear
1	Linear Regression (FE)	13.295618	9.741798	0.759762	Linear
2	Ridge Regression (Raw)	13.320041	9.764298	0.758879	Linear
3	Ridge Regression (FE)	13.295342	9.741233	0.759772	Linear
4	Lasso Regression (Raw)	13.319966	9.764137	0.758882	Linear
5	Lasso Regression (FE)	13.295254	9.741069	0.759775	Linear
6	MLP with SGD	NaN	NaN	NaN	Neural Net
7	MLP with Adam	14.812958	10.779123	0.701800	Neural Net
8	MLP with RMSprop	13.307721	9.686967	0.759325	Neural Net
9	MLP with Adagrad	13.210906	9.647330	0.762814	Neural Net
10	Ensemble of 5 MLPs (Adagrad)	13.271379	9.695907	0.760637	Ensemble
11	Final Optimized MLP (Optuna Hyperparams)	13.203998	9.628222	0.763062	Neural Net
12	Final Optimized LightGBM (Optuna Hyperparams)	12.931977	9.379242	0.772724	Tree-Based
13	Stacked Ensemble (All Models)	12.930729	9.381626	0.772768	Ensemble
14	Final Optimized CatBoost (Optuna Hyperparams)	12.997247	9.454901	0.770424	Tree-Based
15	LightGBM (Final - Kaggle Submission)	12.855765	9.337937	0.775395	Tree-Based
16	XGBoost	12.932367	9.371496	0.772710	Tree-Based
17	Tree-Based Models Ensemble (LGBM + XGB + CatBo...	12.695807	9.247818	0.780949	Ensemble

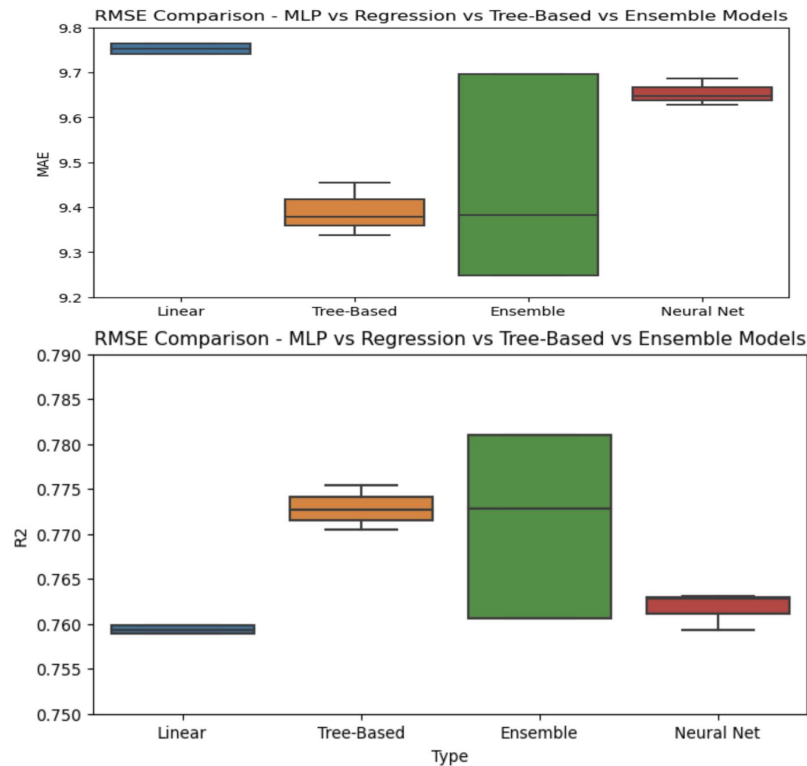
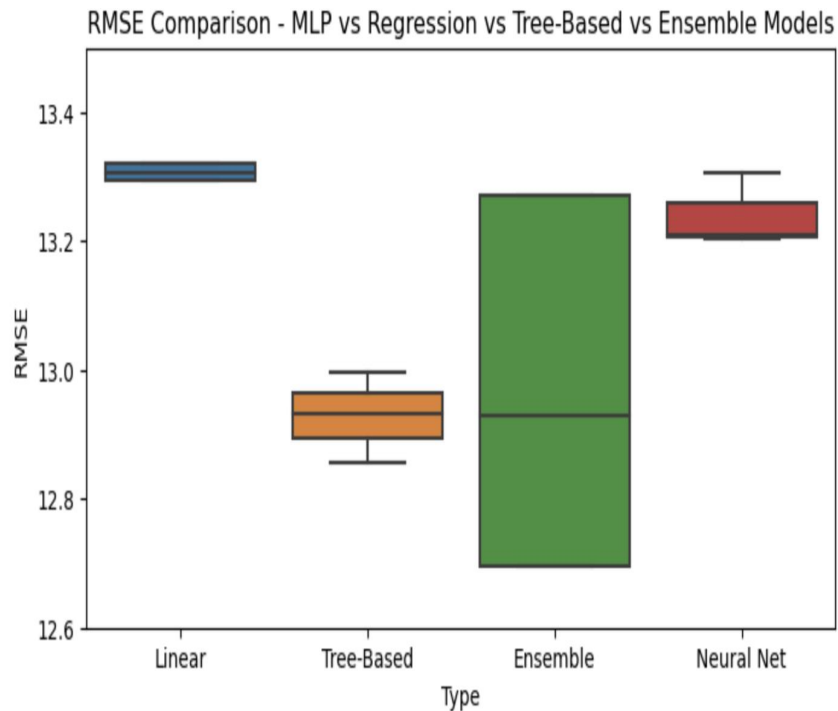
Data Visualizations: Comparison Between Model Types

Bar Charts



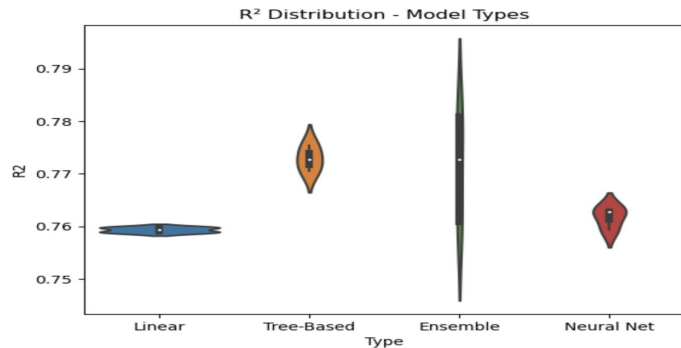


Data Visualizations: Comparison Between Model Types (Cont'd)

Box Charts



Data Visualizations: Comparison Between Model Types (Cont'd)

Violin Charts




Data Visualizations: Comparison Between Model Types (Cont'd)

Outlier!

Parallel Coordinates

Normalized Parallel Coordinates for Model Performance



Data Visualizations: Comparison Between Model Types (Cont'd)

- Overall, the tree based models and the ensembles performed the best
 - Lowest RMSE/MAE scores and highest R^2 score
- Ensemble methods had the highest distribution among the data, and linear models had the lowest distribution among the data
- Parallel coordinate chart reveals an inverse relationship between the mean squared error (RMSE/MAE) metrics and R^2 score
- Higher variance among the ensemble methods could imply sensitivity to outliers, possibly explaining why they performed slightly worse than the tree-based models

Conclusion

- LightGBM proved to be a strong model
 - Efficiency, support for categorical features, robustness to missing data - great for the dataset
 - Feature engineering had significant positive impact on performance
- Future improvements:
 - Cross validation
 - Deeper hyperparameter optimization
 - Ensemble with other models such as XGBoost or CatBoost